

AI GUARDIAN · CYBERSECURITY × FINTECH × HEALTH

AEGIS

"Stop scams before they happen."

An AI guardian that detects manipulation in any conversation — voice, SMS, or chat — in real time, and intervenes *before* the money is lost. One engine, every channel and language.

CYBERSECURITY	FINTECH	HEALTH
Social-engineering & deepfake intent	Real-time transaction shield	Protecting the vulnerable

SUBMISSION

Aegis — The AI Guardian That Stops Scams Before They Happen

LIVE DEMO

[beyond-tomorrow-summit-ai-future-te.vercel.app](#)

STATUS

Working end-to-end prototype · 61/61 tests pass · deployed

ENTRANT

Tumo Mogame — solo (full-stack, AI, design, pitch)

REPOSITORY

[github.com/TUMO-MOGAME/Beyond-Tomorrow-Summit-AI-Future-Tech-Ecosystem-](#)

CONTACT

tumomogame9@gmail.com

01 — THE PROBLEM

AI made scams infinitely scalable. Our defenses didn't keep up.

Generative AI made fraud faster, cheaper, and nearly impossible to detect — and the people who suffer most are those least equipped to fight back.

\$40B Projected US AI-enabled fraud losses by 2027 (up from \$12.3B in 2023 — 32% CAGR). DELOITTE	3 sec Of audio is enough to clone a voice at 85% accuracy. 70% of people can't tell it's fake. AARP
77% Of voice-clone scam victims lose money; 25% of adults have been affected or know someone who has. MCAFEE	5 min Average interval between deepfake fraud attempts — and accelerating. ENTRUST

WHY TODAY'S DEFENSES FAIL


01	Reactive	Banks and existing tools flag fraud <i>after</i> the transfer clears. By then the money — often irreversible — is already gone.
02	Keyword-based	Blocklist and keyword filters are trivially evaded — AI-generated scams rewrite themselves on every attempt and walk straight past them.
03	No human moment	There is no intervention while the victim is being manipulated. They are alone with the attacker at the one moment that matters.

Everything today is forensics. The problem demands prevention.

02 — HOW IT WORKS

A real-time guardian that detects intent — and intervenes.

The Scam DNA Engine analyzes any conversation and detects the underlying manipulation structure — urgency, authority impersonation, secrecy, fear, irreversible payment — not surface keywords. When risk crosses the threshold on a money scam, Aegis acts.



- 1 Ingest** ANY CHANNEL
Paste or simulate a live call, SMS, or chat. One pipeline, modality-agnostic.
- 2 Scam DNA Engine** LLM · ZOD
An LLM returns structured JSON: risk score, the exact triggering quotes, the manipulation tactics, a plain-language summary, and a recommended action.
- 3 Warn + alert the Trusted Circle** HUMAN-IN-LOOP
The user sees why it's dangerous in plain language; a pre-set family member is alerted — the most effective real-world circuit-breaker for elder fraud.
- 4 Freeze the payment — Transaction Shield** FINTECH
The suspicious transfer is held before it completes, stopping the loss at the moment it matters.

WHY AEGIS IS DIFFERENT — THE MOAT

01 Prevention, not forensics Acts <i>during</i> the attack, before the payment leaves — not after the money is gone.	02 Intent, not keywords One engine generalizes across every channel and language. The Spanish demo proves it — no retraining.	03 Trusted Circle + Care Mode Loops in family. The senior does nothing — family sets it up once. A social safety net, not a lonely alert.
---	--	--

03 — SEE IT IN 30 SECONDS

Run the rescue.

- 01 Open the live demo and go to the **Guardian** dashboard.
- 02 Hit **Simulate call** on the grandparent fixture — the scam transcript streams in.
- 03 Watch the **risk meter climb** and the manipulation tactics light up with the exact words that triggered them.
- 04 At threshold, the **Trusted Circle alert** fires and the **Transaction Shield** freezes the transfer.
- 05 Run the **Spanish** fixture to prove intent-detection is language-agnostic; run the **safe** fixture to see a benign chat score low.

BUILD STATUS

● 61 / 61 AUTOMATED CHECKS PASS

Prototype	End-to-end, live
Deploy	Vercel (public URL)
Repo	Public on GitHub
Fail-safe	Mock engine — demo never hard-fails

PROVIDER-AGNOSTIC ENGINE

Auto-detects an API key — Anthropic → Groq → Gemini → OpenRouter → OpenAI — and falls back to an offline heuristic **MOCK** if none is present. Every integration degrades gracefully without a key.

TECHNOLOGY STACK

AI / ML Claude (Anthropic) LLM · structured tool-based JSON · prompt caching · Scam DNA system prompt · speech-to-text (Whisper / AssemblyAI)

FRONTEND Next.js 14 (App Router) · TypeScript · Tailwind · Framer Motion **BACKEND** Next.js API routes · Node.js · Zod validation-with-retry

INTEGRATIONS Resend (email) · Twilio (SMS / Guardian Number call-screening) **INFRA** Vercel · GitHub

WHAT'S NEXT

NEAR TERM	Dedicated synthetic-voice (deepfake-audio) forensic model as a second signal.
DISTRIBUTION	B2B2C with banks, telcos & insurers; an SDK so any banking or messaging app can embed Aegis.
PRIVACY	On-device / edge inference — conversations never leave the user's phone.
REACH	Expanded multilingual coverage and eldercare / caregiver-dashboard integrations.

"Scammers got an AI. It's time the rest of us got one too. We turn a \$40-billion problem into a moment of safety."

TUMO MOGAME · AEGIS · TUMOMOGAME9@GMAIL.COM · GITHUB.COM/TUMO-MOGAME